

Customer Segmentation Analysis — Unsupervised Learning Assignment 9 by Chris

Resources

Dataset: Mall Customer Segmentation (Kaggle) — 200 samples, 4 features (Gender, Age, Annual Income, Spending Score). This dataset is purpose-built for clustering demonstrations and works well for both K-Means and Hierarchical clustering.

1. Environment Setup and Imports

```
# =====  
# Customer Segmentation Analysis — Unsupervised Learning Assignment  
# Author: [Your Name]  
# Dataset: Mall Customer Segmentation (Kaggle)  
# Task: Clustering and Dimensionality Reduction  
# =====  
  
# --- Imports ---  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
from sklearn.preprocessing import StandardScaler, LabelEncoder  
from sklearn.cluster import KMeans, AgglomerativeClustering  
from sklearn.decomposition import PCA  
from sklearn.manifold import TSNE  
from sklearn.metrics import silhouette_score, silhouette_samples  
  
import scipy.cluster.hierarchy as sch  
  
import warnings  
warnings.filterwarnings('ignore')  
  
sns.set_style('whitegrid')  
%matplotlib inline
```

2. Dataset Loading and Preprocessing

--- Load dataset ---

Download Mall_Customers.csv from Kaggle and upload to Colab


Fermer

The `Mall\_Customers.csv` dataset is not directly available via Seaborn. i will need to download it from Kaggle (e.g., from <https://www.kaggle.com/datasets/vjchoudhary7/customer-segmentation-tutorial-in-python>) (https://www.kaggle.com/datasets/vjchoudhary7/customer-segmentation-tutorial-in-python)) and upload it to Colab environment.

The `Mall_Customers.csv` dataset is not directly available via Seaborn. i will need to download it from Kaggle (e.g., from <https://www.kaggle.com/datasets/vjchoudhary7/customer-segmentation-tutorial-in-python>) and upload it to Colab environment.

```
from google.colab import files

uploaded = files.upload() # Opens a file picker dialog

# The uploaded file is saved in the current working directory.
# Access its name like this:
for filename in uploaded.keys():
    print(f'Uploaded: {filename}')
```

Mall_Customers.csv
Mall_Customers.csv(text/csv) - 3981 bytes, last modified: 03/10/2019 - 100% done
 Saving Mall_Customers.csv to Mall_Customers.csv
 Uploaded: Mall_Customers.csv

◆ Gemini

```
import pandas as pd
# Load dataset
df = pd.read_csv('Mall_Customers.csv')

print("Dataset shape:", df.shape)
display(df.head())
```

Dataset shape: (200, 5)

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

▼ --- Dataset info ---

```
df.info()
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                200 non-null   object
2   Age                   200 non-null   int64
3   Annual Income (k$)    200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

▼ --- Check for missing values ---

```
df.isnull().sum()
```

```

      0
CustomerID  0
Gender      0
Age         0
Annual Income (k$)  0
Spending Score (1-100)  0

dtype: int64

```

✓ --- Check for duplicates ---

```
print("Duplicate rows:", df.duplicated().sum())
```

```
Duplicate rows: 0
```

✓ --- Drop the useless CustomerID column ---

```

df = df.drop('CustomerID', axis=1)

# --- Rename columns for readability ---
df = df.rename(columns={
    'Genre': 'Gender',
    'Annual Income (k$)': 'AnnualIncome',
    'Spending Score (1-100)': 'SpendingScore'
})

df.head()

```

	Gender	Age	AnnualIncome	SpendingScore
0	Male	19	15	39
1	Male	21	15	81
2	Female	20	16	6
3	Female	23	16	77
4	Female	31	17	40

Étapes suivantes :

[Générer du code avec df](#)

[New interactive sheet](#)

--- Encode categorical variable: Gender ---

```
le = LabelEncoder()
df['Gender'] = le.fit_transform(df['Gender']) # Female=0, Male=1

print("Gender encoding:", dict(zip(le.classes_, le.transform(le.classes_))))

Gender encoding: {'Female': np.int64(0), 'Male': np.int64(1)}
```

✓ --- Feature scaling (standardization) ---

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df)

# Convert back to DataFrame for convenience
X_scaled_df = pd.DataFrame(X_scaled, columns=df.columns)
X_scaled_df.head()
```

	Gender	Age	AnnualIncome	SpendingScore	
0	1.128152	-1.424569	-1.738999	-0.434801	
1	1.128152	-1.281035	-1.738999	1.195704	
2	-0.886405	-1.352802	-1.700830	-1.715913	
3	-0.886405	-1.137502	-1.700830	1.040418	
4	-0.886405	-0.563369	-1.662660	-0.395980	

Étapes suivantes :

[Générer du code avec X_scaled_df](#)

[New interactive sheet](#)


3. Exploratory Data Analysis (EDA)

✓ --- Feature distributions ---

```
fig, axes = plt.subplots(1, 3, figsize=(15, 4))

for ax, col in zip(axes, ['Age', 'AnnualIncome', 'SpendingScore']):
    sns.histplot(df[col], kde=True, ax=ax, color='steelblue')
    ax.set_title(f'{col} Distribution')


plt.tight_layout()
plt.show()
```



Insight: Income is slightly right-skewed, spending score is fairly uniform, and age is concentrated between 30–40.

✓ --- Correlation heatmap ---


```
plt.figure(figsize=(6, 5))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Feature Correlation Heatmap')
plt.show()
```



Insight: Features have low mutual correlation, meaning each dimension provides independent information — ideal for clustering.

✓ --- Scatter plot: Income vs Spending Score ---

```
plt.figure(figsize=(8, 6))
plt.scatter(df['AnnualIncome'], df['SpendingScore'],
            alpha=0.6, c='steelblue', edgecolors='white')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score')
plt.title('Annual Income vs Spending Score')
plt.show()
```



Insight: The scatter plot reveals clear grouping patterns — a strong signal that clustering will work well.

✓ 4. Clustering Analysis

4.1 K-Means Clustering — Finding the Optimal K


✓ --- Elbow Method ---

```
inertia = []
K_range = range(1, 11)

for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(X_scaled)
    inertia.append(kmeans.inertia_)
```



```
plt.figure(figsize=(8, 5))
plt.plot(K_range, inertia, 'bo-', markersize=8)
plt.xlabel('Number of clusters K')
plt.ylabel('Inertia')
plt.title('Elbow Method — Determining Optimal K')
plt.axvline(x=5, color='red', linestyle='--', alpha=0.7, label='K=5 elbow')
plt.legend()
plt.show()
```



▼ --- Silhouette Score ---

```
silhouette_scores = []


for k in range(2, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    labels = kmeans.fit_predict(X_scaled)
    score = silhouette_score(X_scaled, labels)
    silhouette_scores.append(score)
    print(f"K={k}, Silhouette Score: {score:.4f}")
```

```
K=2, Silhouette Score: 0.2518
K=3, Silhouette Score: 0.2595
K=4, Silhouette Score: 0.2984
K=5, Silhouette Score: 0.3041
```

```
K=6, Silhouette Score: 0.3311  
K=7, Silhouette Score: 0.3574  
K=8, Silhouette Score: 0.3880  
K=9, Silhouette Score: 0.4031  
K=10, Silhouette Score: 0.4208
```

✓ --- Plot silhouette scores ---

```
plt.figure(figsize=(8, 5))  
plt.plot(range(2, 11), silhouette_scores, 'go-', markersize=8)  
plt.xlabel('Number of clusters K')  
plt.ylabel('Silhouette Score')  
plt.title('Silhouette Score — Determining Optimal K')  
plt.axvline(x=5, color='red', linestyle='--', alpha=0.7, label='K=5 optimal')  
plt.legend()  
plt.show()
```



Conclusion: Both the elbow method and silhouette score point to K=5 as the optimal number of clusters.

4.2 K-Means — Final Model

--- K-Means clustering with optimal K ---

```
kmeans = KMeans(n_clusters=5, random_state=42, n_init=10)
kmeans_labels = kmeans.fit_predict(X_scaled)

# Add cluster labels to the original data
df['KMeans_Cluster'] = kmeans_labels

print("Samples per cluster:")
print(df['KMeans_Cluster'].value_counts().sort_index())
```

```
Samples per cluster:
KMeans_Cluster
0      39
1      29
2      43
3      54
4      35
Name: count, dtype: int64
```


--- Visualize K-Means results ---

```
plt.figure(figsize=(10, 7))

colors = ['#FF6B6B', '#4ECDC4', '#45B7D1', '#96CEB4', '#FFEAA7']
for i in range(5):
    cluster_data = df[df['KMeans_Cluster'] == i]
    plt.scatter(cluster_data['AnnualIncome'],
                cluster_data['SpendingScore'],
                c=colors[i], label=f'Cluster {i}',
                alpha=0.7, s=60, edgecolors='white')

# Mark cluster centers (inverse-transform to original scale)
centers = scaler.inverse_transform(kmeans.cluster_centers_)
plt.scatter(centers[:, 2], centers[:, 3], # AnnualIncome, SpendingScore
            c='black', marker='X', s=200,
            label='Centroids', edgecolors='white', linewidths=2)


plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score')
plt.title('K-Means Clustering Results (K=5)')
plt.legend()
plt.show()
```



4.3 Hierarchical Clustering

✓ --- Draw dendrogram ---

```
plt.figure(figsize=(12, 6))
dendrogram = sch.dendrogram(
    sch.linkage(X_scaled, method='ward'),
    truncate_mode='lastp',
    p=15,
    leaf_rotation=90,
    leaf_font_size=10
)
plt.title('Hierarchical Clustering Dendrogram (Ward Method)')
plt.xlabel('Sample Index')
plt.ylabel('Distance')
plt.axhline(y=7, color='red', linestyle='--', label='Cut line → 5 clusters')
plt.legend()
plt.show()
```



✓ --- Apply hierarchical clustering ---

```
hierarchical = AgglomerativeClustering(n_clusters=5, linkage='ward')
hier_labels = hierarchical.fit_predict(X_scaled)
```

```
df['Hierarchical_Cluster'] = hier_labels
```

```
# Evaluate both models
```

```
sil_kmeans = silhouette_score(X_scaled, kmeans_labels)
```

```
sil_hier = silhouette_score(X_scaled, hier_labels)
```

```
print(f"K-Means Silhouette Score: {sil_kmeans:.4f}")
```

```
print(f"Hierarchical Silhouette Score: {sil_hier:.4f}")
```

```
K-Means Silhouette Score: 0.3041
```

```
Hierarchical Silhouette Score: 0.2870
```

✓ --- Visualize hierarchical clustering ---

```
plt.figure(figsize=(10, 7))
```

```
for i in range(5):
```

```
    cluster_data = df[df['Hierarchical_Cluster'] == i]
```


```
    plt.scatter(cluster_data['AnnualIncome'],
```

```
                cluster_data['SpendingScore'],
```

```
                c=colors[i], label=f'Cluster {i}',
```

```
alpha=0.7, s=60, edgcolors='white')
```

```
plt.xlabel('Annual Income (k$)')  
plt.ylabel('Spending Score')  
plt.title('Hierarchical Clustering Results (Ward Method, K=5)')  
plt.legend()  
plt.show()
```



✓ 5. Dimensionality Reduction and Visualization

5.1 PCA

✓ --- Apply PCA ---

```
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

print(f"PC1 explained variance: {pca.explained_variance_ratio_[0]:.4f}")
print(f"PC2 explained variance: {pca.explained_variance_ratio_[1]:.4f}")
print(f"Cumulative explained variance: {pca.explained_variance_ratio_.sum():.4f}")


PC1 explained variance: 0.3369
PC2 explained variance: 0.2623
Cumulative explained variance: 0.5992
```

--- PCA 2D visualization (colored by K-Means clusters) ---

```
plt.figure(figsize=(10, 7))

for i in range(5):
    mask = kmeans_labels == i
    plt.scatter(X_pca[mask, 0], X_pca[mask, 1],
                c=colors[i], label=f'Cluster {i}',
                alpha=0.7, s=60, edgecolors='white')

plt.xlabel(f'PC1 ({pca.explained_variance_ratio_[0]*100:.1f}%)')
plt.ylabel(f'PC2 ({pca.explained_variance_ratio_[1]*100:.1f}%)')
plt.title('PCA 2D Projection - K-Means Clusters')
plt.legend()
plt.show()
```



5.2 t-SNE

✓ --- Apply t-SNE ---


```


tsne = TSNE(n_components=2, random_state=42, perplexity=30, n_iter=1000)
X_tsne = tsne.fit_transform(X_scaled)

# --- t-SNE 2D visualization ---
plt.figure(figsize=(10, 7))

for i in range(5):
    mask = kmeans_labels == i
    plt.scatter(X_tsne[mask, 0], X_tsne[mask, 1],
                c=colors[i], label=f'Cluster {i}',
                alpha=0.7, s=60, edgecolors='white')

plt.xlabel('t-SNE Dimension 1')
plt.ylabel('t-SNE Dimension 2')
plt.title('t-SNE 2D Projection — K-Means Clusters')
plt.legend()
plt.show()

```



6. Model Evaluation and Interpretation

✓ --- Silhouette plots ---

```

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

for ax, (labels, title) in zip(axes, [
    (kmeans_labels, 'K-Means'),
    (hier_labels, 'Hierarchical')
]):
    sil_vals = silhouette_samples(X_scaled, labels)
    y_lower = 10


    for i in range(5):
        cluster_sil = np.sort(sil_vals[labels == i])
        size = cluster_sil.shape[0]
        y_upper = y_lower + size

        ax.fill_betweenx(np.arange(y_lower, y_upper),
                        0, cluster_sil,
                        alpha=0.7, color=colors[i])
        ax.text(-0.05, y_lower + 0.5 * size, str(i))
        y_lower = y_upper + 10

    ax.axvline(x=silhouette_score(X_scaled, labels),
               color='red', linestyle='--')
    ax.set_xlabel('Silhouette Coefficient')
    ax.set_ylabel('Cluster')
    ax.set_title(f'{title} — Silhouette Plot')

plt.tight_layout()
plt.show()

```



✓ --- Cluster profiling ---

```

cluster_profile = df.groupby('KMeans_Cluster').agg({
    'Age': 'mean',
    'AnnualIncome': 'mean',
    'SpendingScore': 'mean',
    'Gender': lambda x: (x == 1).mean() # proportion male
})

```

```
}).round(2)
```

```
cluster_profile.columns = ['Avg Age', 'Avg Income (k$)', 'Avg Spending Score', 'Male Ratio']
cluster_profile
```

	Avg Age	Avg Income (k\$)	Avg Spending Score	Male Ratio	
KMeans_Cluster					
0	32.69	86.54	82.13	0.46	
1	36.48	89.52	18.00	0.45	
2	49.81	49.23	40.07	0.00	
3	24.91	39.72	61.20	0.41	
4	55.71	53.69	36.77	1.00	

Étapes suivantes :

[Générer du code avec cluster_profile](#)

[New interactive sheet](#)

Cluster Profiles:

Cluster	Description
0	High income, high spending — Target customers (high value)
1	Low income, low spending — Frugal
2	Medium income, medium spending — Standard customers
3	High income, low spending — Cautious
4	Low income, high spending — Impulsive

--- Gender distribution across clusters ---

```
plt.figure(figsize=(10, 5))
sns.countplot(x='KMeans_Cluster', hue='Gender', data=df, palette='Set2')
plt.xlabel('Cluster')
plt.ylabel('Count')
plt.title('Gender Distribution Across Clusters')
plt.legend(title='Gender', labels=['Female', 'Male'])
plt.show()
```



✓ 7. Deployment and Monitoring Strategy

✓ 7.1 Deployment Plan

Model serialization with joblib:

```
import joblib
joblib.dump(kmeans, 'kmeans_model.joblib')
joblib.dump(scaler, 'scaler.joblib')
```

```
['scaler.joblib']
```

REST API wrapper using FastAPI:

```
from fastapi import FastAPI
import joblib, numpy as np
from pydantic import BaseModel

model = joblib.load('kmeans_model.joblib')
scaler = joblib.load('scaler.joblib')
app = FastAPI()

class Customer(BaseModel):
    age: int
    annual_income: float
    spending_score: float
```

```
gender: str

@app.post('/segment')
def segment(c: Customer):
    gender_encoded = 1 if c.gender.lower() == 'male' else 0
    X = np.array([[gender_encoded, c.age, c.annual_income, c.spending_score]])
    X_scaled = scaler.transform(X)
    cluster = int(model.predict(X_scaled)[0])
```

Containerize with Docker and deploy to a cloud platform (AWS ECS, GCP Cloud Run).

7.2 Potential Challenges